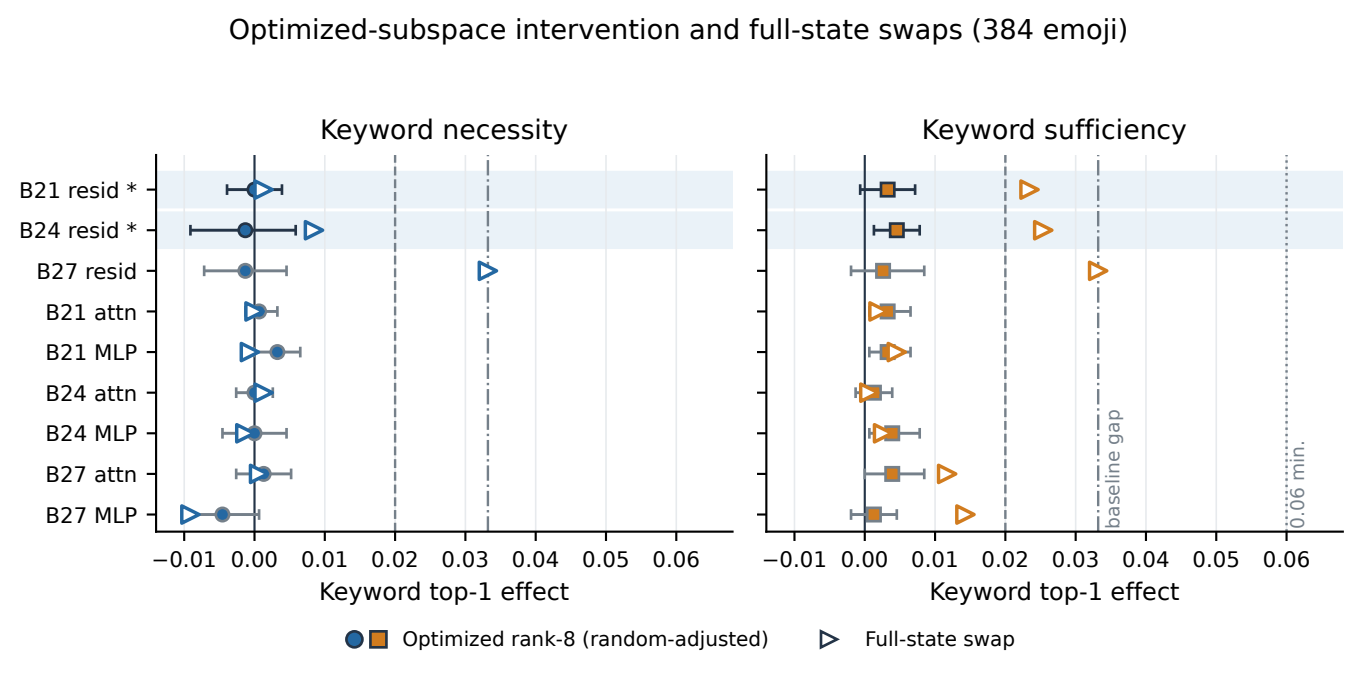
Optimized-subspace intervention and full-state swaps (384 emoji)
Keyword necessity
Keyword sufficiency
B21 resid *
B24 resid *
B27 resid
B21 attn
B21 MLP
B24 attn
B24 MLP
B27 attn
B27 MLP
−0.01
0.00
0.01
0.02
0.03
0.04
0.05
0.06
Keyword top-1 effect
baseline gap
0.06 min.
Optimized rank-8 (random-adjusted)
Full-state swap